BUSINESS CASE · ILLUSTRATIVE ASSUMPTIONS

A focused planning product with an implementation-led path

Commercial hypotheses for validation with resilience teams; not a forecast, offer, or customer quote.

Customer value hypothesis

ContinuityMap could reduce repeated manual assembly of dependency-aware exercises, make evidence gaps visible earlier, improve coordination between service owners, and produce traceable review records. These benefits are not quantified by the synthetic application. A pilot should measure preparation time, graph coverage, owner agreement with impact paths, and follow-up action closure.

Illustrative delivery and pricing models

Possible offers include an annual organization subscription, a fee for each site, a customer-hosted license, setup and data-mapping work, optional scenario modules, and professional assessments. This example assumes a $65,000 annual platform fee plus $9,000 per site for three sites, or $92,000 recurring per organization. First-year revenue adds $28,000 onboarding.

Organizations	Year-one revenue	Direct cost	Gross profit	Gross margin	Annual subscription total
1	$120,000	$30,000	$90,000	75%	$92,000
5	$600,000	$150,000	$450,000	75%	$460,000
10	$1,200,000	$300,000	$900,000	75%	$920,000
25	$3,000,000	$750,000	$2,250,000	75%	$2,300,000

How to read this example: direct cost is the assumed cost to onboard and serve an organization. Gross profit = revenue – direct costs. Gross margin = gross profit ÷ revenue. For one organization, $120,000 – $30,000 = $90,000 gross profit; $90,000 ÷ $120,000 = 75%. The $30,000 assumes $18,000 onboarding delivery plus $12,000 yearly hosting/support. Year two would have $80,000 gross profit and 87% gross margin only if revenue stays $92,000 and direct cost stays $12,000. These figures exclude product development, sales, management, taxes, security assurance, and other company overhead. They are assumptions, not a forecast.

VALUE SENSITIVITY · NOT A PROMISE

Labor savings alone may not support the entry price

A transparent sensitivity range exposes the willingness-to-pay question rather than hiding it.

Example	Staff hours saved / exercise	Exercises / year	Staff cost / hour	Possible yearly staff-time value
Low	40	3	$60	$7,200
Middle	80	4	$100	$32,000
High	160	6	$140	$134,400

Formula: staff hours saved per exercise × exercises per year × staff cost per hour. No target customer has validated these inputs. The middle example is less than the assumed $92,000 annual subscription, so staff-time savings alone may not justify the price for a single-site or low-frequency buyer.

The commercial case may depend on multi-site rollups, reduced consultant effort, audit traceability, repeatable exercises, service-owner coordination, and lower cost of updating plans. Any claimed cash value requires customer baseline data and a controlled before/after study. The product cannot guarantee avoided outages or improved regulatory status.

Customer segments and purchase constraints

Segment	Potential job	Adoption issue
Municipal resilience office	Cross-agency exercise preparation	Data-sharing authority and owner coordination
Power / water utility	Facility and service dependency review	Sensitive operational data and system integration
Hospital / airport	Continuity scenario traceability	Procurement, identity, and existing BCM stack
Consultancy / government	Repeatable assessment and on-prem delivery	Custom data onboarding and assurance burden

COSTS · RISKS · GO-TO-MARKET

Customer data quality is the main delivery risk

Commercial success depends on repeatable onboarding and customer trust in the evidence model.

Cost / risk	Why it matters	Mitigation hypothesis
Dependency data collection	Relationships may be scattered, stale, or contested.	Customer-specific data mapping, owner workshops, and validation status.
Integration effort	Identifiers, plans, and access controls vary by customer.	Start with file imports; price implementation separately.
Sensitive data	Infrastructure context may be restricted or security-sensitive.	Customer-hosted/on-prem options; minimize retained document text.
AI trust	Extraction and summaries can be wrong or misleading.	Citations, review-only proposals, local inference, owner sign-off.
Incumbent overlap	Existing platforms may already map services and dependencies.	Compare workflow fit, exportability, deployment, and evidence control.
Long procurement	Government and utilities can have lengthy approvals.	Pilot with a bounded, synthetic or approved data subset.

Go-to-market sequence

- 0–3 months: 8–12 discovery conversations, data-boundary review, and one safe exercise use case.
- 3–6 months: customer-controlled pilot with named owners, one critical service cohort, and baseline preparation metrics.
- 6–12 months: harden identity and audit, prove implementation repeatability, and compare against incumbent tools.

The key commercial gate is whether customers can produce a current, evidence-backed graph at a repeatable cost and whether the workflow offers value beyond existing continuity and event-management platforms. See docs/SOURCES.md for research and product sources.